

Standardized Type Ordering

Document #: P2830R2
Date: 2024-03-10
Project: Programming Language C++
Audience: EWG
Reply-to: Nate Nichols
<natenichols@cox.net>
Gašper Ažman
<gasper.azman@gmail.com>

Contents

1	Abstract	2
2	Revision History	2
3	Motivation	2
3.1	Motivating Examples	3
3.1.1	Canonicalizing policy-based libraries	3
3.1.2	Canonicalized variant	4
3.1.3	Canonicalized tuple	5
3.1.4	API stability	6
4	Design discussion	6
4.1	Why implementation-defined?	6
4.1.1	ABI specifications already have to do this work	6
4.1.2	The C++ standard doesn't own cross-compilation-unit semantics	6
4.1.3	Any new proposal to C++ would have to consider the ordering	7
4.1.4	ORDER(X, Y) already has public-facing API implications	7
4.2	Desirable properties of ORDER(x, y)	7
4.2.1	Stability	7
4.2.2	Free-standing	7
4.2.3	Consistency with <code>type_info::before()</code>	7
4.2.4	Self-consistency	7
4.2.5	Reflection compatibility	7
5	Proposal	8
5.1	Semantics	8
5.2	Syntax	8
5.2.1	Option 1 (preferred): a variable template <code>std::type_order_v<T, U></code>	8
5.2.2	Option 2: variable template <code>std::entity_ordering<X, Y></code>	9
5.2.3	Option 3: reflection	9
5.2.4	Option 4: heterogeneous <code>constexpr std::type_identity::operator<=></code>	9
5.2.5	Option 5: <code>constexpr std::_lift<arg>::operator<=></code>	10
5.2.6	Non-Option: <code>constexpr bool std::type_info::before()</code>	10
6	FAQ	10
6.1	Why should this be standardized?	10
6.2	Why not wait for reflection?	10
6.3	But couldn't this be done faster with reflection?	10

7	Acknowledgements	10
8	Proposed Wording	11
9	Full order proposal (former alternative)	11
9.1	Approach	11
9.2	Structure of key-tuples	11
9.3	Named Scopes	11
9.3.1	Example 1: <code>class foo</code> is declared in <code>struct bar</code> :	12
9.3.2	Example:	12
9.4	Atoms	12
9.5	Kinds	12
9.5.1	Identifiers	13
9.5.2	Namespaces	14
9.5.3	Types	14
9.5.4	Ordering Scalar Types	15
9.5.5	Ordering Array Types	16
9.5.6	Ordering Class Types	16
10	Appendix A: building <code>apply_canonicalized</code>	20
10.1	Full code listing as tested and implemented	21
11	References	22

1 Abstract

Currently, `std::type_info` provides a stable but *implementation defined* order of types. This paper explores a standardized ordering of types in C++, and possible syntaxes for it.

This paper is split this paper into two parts:

1. the design of the exposition-only function `ORDER(x, y)`, which deals with how such an order should be defined,
2. how to expose this capability to the programmer.

2 Revision History

0. New Paper

1. Revision 1

- Introduce options to prevent changing `std::type_info::before`
- Anonymous namespaces can't be empty
- Add FAQ section
- Add motivating examples
- Add proposed syntax
- Add appendix

2. Revision 2

- Propose to just make the implementation-defined order `constexpr`, as suggested by EWGi and BSI reviewers.

3 Motivation

There is currently no way in portable C++ to sort types at compile-time.

Various libraries exist to do this, mostly by utilizing `__PRETTY_FUNCTION__`, but this is non-portable and error-prone (consider forward-declared enums). There are multiple stack-overflow questions on the topic.

One such implementation is part of the monadic functional library by Bronek Kozicki et al.

Fundamentally, we need a way to canonicalize sets and multisets of types. This is necessary when building any kind of compositional library in C++, from `std::execution` [P2300R7], mixin libraries, libraries of monadic components and operators, policy-based libraries, etc. The inability to sort and unique types leads to the instantiation of an exponential number of templates instead of just one per typeset, which leads to code-bloat and untenable compile-times. The problem is fundamentally that without canonicalization, we generate types with a different name but the same behavior, and there is no way to avoid this.

The goal here is to provide a *flexible* language mechanism to let both `Foo<A, B, C>` and `Foo<C, B, A>` produce the same underlying `Foo_impl<A, B, C>`.

The reason we start with `ORDER()` and not “just give me typesets” is that we need flexibility; consider

- a given library might want to deduplicate on a part and keep either last or first, or even make it ill-formed: `Foo<pair<A, X>, pair<B, Y>, pair<A, Z>>` might want to be the same as `Foo_impl<pair<A, X>, pair<B, Y>>` or `Foo_impl<pair<A, Z>, pair<B, Y>>`
- a given library might actually just want canonicalized multisets: `Foo<A, B, A, A, C>` should perhaps be `Foo<A, A, A, B, C>`
- or treat the first one as special: `Matrix<float, policy1, policy2, policy3>` should only deduplicate policies.

We must provide `ORDER` in order to sort and unique; they are required building blocks for any `set` primitive. Put another way, even if we standardized a `set`, we’d need to somehow canonicalize the order (due to mangling and debug info), leading us back here.

Without such canonicalization, it is also functionally impossible to enumerate the set of function templates to instantiate in a separate compilation unit. For the same reason, it’s utterly impossible to type-erase them in a fixed-size virtual function table.

3.1 Motivating Examples

This section introduces the kind of code we would like to write, regardless of how this feature ends up being spelled. To not prejudice the reader in this section, please assume the existence of an exposition-only `constexpr` function `ORDER(x, y) -> std::strong_ordering` whose arguments can be anything usable as a template argument. Crucially, also consider the interactions with [P1985R3] and [P2841R1], which introduce `constexpr` and `variable-template` arguments.

Note: `ORDER(x, y)` can be defined just on types by ordering a class template that takes `x` and `y` as a template argument, so we can realistically just define it on types.

3.1.1 Canonicalizing policy-based libraries

Consider the needs of a library for type-erasure; we’d like the user to specify the capabilities to erase. Fundamentally, these capabilities are a set.

Observe:

```
using T1 = lib::dyn<ostreamable, json_serializable, iterable<int>, scalar_multiplicable<int>>;
```

Different users of the library will likely provide these capabilities in different orders; this becomes especially problematic when writing functions:

```
using T2 = lib::dyn<json_serializable, ostreamable, iterable<int>, scalar_multiplicable<int>>;
~~~~~ flipped ~~~~~
void f(T2 const&);
```

We can solve this by having type aliases, but if these sets are *computed*, we are left without recourse:

```
int sum = 0;
auto pipeline1 =
    log(std::cerr) | sum_into(&sum) | imul(sum) | json_dump(std::cout);
// dyn<ostreamable, iterable<int>, scalar_multiplicable<int>, json_serializable>
auto pipeline2 =
    sum_into(&sum) | imul(sum) | json_dump(std::cout) | log(std::cerr);
// dyn<iterable<int>, scalar_multiplicable<int>, json_serializable, ostreamable>
```

The above pipelines need the same type-erased interface for its input, and neither is either T1 or T2.

3.1.2 Canonicalized variant

The most obvious example is a canonicalized `std::variant`, that is, something like

```
template <typename... Ts>
using canonical_variant = apply_canonicalized<std::variant, Ts...>;
```

Building `apply_canonicalized` is not terribly difficult, as long as we have `ORDER`. Please see the appendices on how to do it.

It *would* be nice if `apply_canonicalized` was a language built-in, but to do that, we need to first define `ORDER(x, y)`. After we define `ORDER`, putting `apply_canonicalized` into `type_traits` is a far simpler proposition.

Note: `apply_canonicalized` is roughly `mp11::mp_sort + mp11::unique` with the order derived from `ORDER`.

Effectively any time a variant is type-indexed instead of integer-indexed, we would prefer the compiler generate just one type per set, instead of type-per-list.

Examples follow, but they are by no means exhaustive.

3.1.2.1 Multiple error kinds and `std::expected`

Consider a `std::expected` for an algorithm that can fail in several ways:

```
auto read_packet(socket& sock)
    -> expected<message, variant<io_error, decode_error>>;
```

Such an algorithm doesn't care about whether the error type is `variant<io_error, decode_error>` or `variant<decode_error, io_error>`. This is not as much of a problem in cases where the programmer writes the types by hand, but consider the generic situation where several errors must be aggregated generically:

```
template <typename Decoder>
auto read_packet(socket& sock)
    -> expected<Decoder::message_type,
                variant</*uniqued and sorted io_error + Decoder::error_types*/>>;
```

If we had member packs and pack aliases, and a nice built-in `set`, we could do the above as

```
template <typename... Ts>
...using set = __builtin_uniqued_sorted<Ts...>;

variant<...set<io_error, Decoder::...error_types...>...>;
```

3.1.2.2 Suspension-point storage for asynchronous code

`std::execution` (and every other async library) needs a variant-type to store the current result at a suspension point. In general, since it models a “suspended” function call, it's analogous to some kind of

```
variant<
  closure_tuple<set_value, Arg1_1, Arg1_2, Arg1_3>, // first possible entry point
  closure_tuple<set_value, Arg2_1>,                // second possible entry point
  closure_tuple<set_error, ...>,
  ...
>
```

We call it a **notification**, as a nod to `rxcpp`.

Consider the `transfer(scheduler)` algorithm; it receives all `set_value(PACK)`, `set_error(PACK)` and possibly `set_stopped()` parameter lists, *stores them in a* `variant<tuple<CHANNEL, PACK>...>`, and resumes with the values on a different scheduler once the compute resource becomes available.

The code using it looks like this:

```
some_computation | transfer(scheduler) | then([](auto&&...args){/*...*/})...
```

All the possible closure tuples are fundamentally unsorted, and generating the assembly, debug information, and everything else for this type more than once is a waste of compile-time and executable space (as well as instruction cache).

Furthermore, stamping out this type multiple times also duplicates the calling code, and given the different indices, COMDAT folding cannot deduplicate those code-paths.

3.1.2.3 Compositional State machine models

The states of most state-machines are fundamentally unordered, as are the message types they receive. If a state machine is generated (such as in a compile-time regex library), canonicalizing identically-behaving components would lead to smaller state-machines and faster compile- and execution times.

3.1.3 Canonicalized tuple

Similarly to a canonicalized `variant`, a `canonicalized_tuple` could be implemented as

```
template <typename...Ts>
using canonical_tuple = apply_canonicalized<tuple, Ts...>;
// or
using canonical_tuple = tuple<...set<Ts...>...>;
```

This is far less useful on its own than the `variant`, since initializing that type will be quite the chore. Instead, canonicalized tuples appear as a result of the `environment` monads.

`std::execution` mixes in an environment, which consists at least of a `stoppable_token`, a `scheduler`, and possibly other things like a `domain`, and, at least in the authors' implementation, arbitrary other things. An `allocator` is also optionally part of the environment.

The `environment<pair<Tags, Ts>...>` is a product type that is a map from `Tag` to a value of type `T`, say `std::stop_token` to `never_stop_token{}` and `scheduler` to `thread_pool_scheduler`.

One can push and pop things off the environment in different parts of the pipeline; if the pushes come in the wrong order, the environment is difficult to keep canonicalized.

```
auto ss = std::stop_source();
auto tp = std::static_thread_pool(5);
auto env = empty_env{}
  | push_env<stop_token>(ss.get_token())
  | push_env<scheduler_t>(tp.get_scheduler())
  | push_env<mylib::numa_domain_t>(2);
```

If you imagine those three calls to happen in different algorithm transformers, it's quite likely they'll be in a different order, manifesting a different type, which doesn't make much sense - the type behaves identically, this just leads to code bloat.

3.1.4 API stability

The authors have observed a necessity to sometimes split parts of such pipeline compositions into different compilation units due to excessive compile times and repetitive compilation.

This has proven difficult due to the brittle nature of type ordering of the names of explicit template specializations involved. By far, the main culprit has been the lack of a sensible canonicalization of names, as the order changes far more frequently than the set of types.

Consider:

```
do_something()
| error_if_result_odd()      // there is no semantic change
| error_if_we_too_too_long() // if we flip these lines
| ...
```

Flipping the lines will flip the two exit error types in the error list, however.

Having a defined, stable order of types proved invaluable (even if we did hack the type order manually).

4 Design discussion

The first two revisions of this proposal tried to construct a full type ordering from first principles. This is still included as the alternative proposal, should EWG want us to go that way.

4.1 Why implementation-defined?

The overwhelming response of people working on implementations, as well as people sitting on the Itanium ABI standards committee, was that this is a fool's errand, and we need only look at the bugreports for the Itanium ABI to see why.

4.1.1 ABI specifications already have to do this work

As a taster, consider the following conundrum:

```
using what_type = decltype(
    [](this auto&& self,
        decltype([](decltype(self)&){}) x = {}){ return x; }());
// ^^ outer ^^ inner
```

The inner lambda's "name" is clearly dependent on the outer's, but it also goes the other way! The ABI standards already have to deal with such mangling conundrums, and duplicating their work in the C++ standard itself seems highly counterproductive, especially since the ABI standards already accomplish all the stability guarantees we would want.

4.1.2 The C++ standard doesn't own cross-compilation-unit semantics

Another problem is that within the C++ standard, we cannot legislate what happens with orderings between compilation units. Granted, we do not need this for `constexpr`, but we do need it to be stable and make sense, so a certain amount of common sense will be expected from implementations *anyway*.

The recommendation was overwhelmingly to just let the order be implementation-defined, and let the implementations do the sensible thing.

4.1.3 Any new proposal to C++ would have to consider the ordering

“punting it off to implementations” saves the C++ standardization process from having to figure this out for every new proposal that touches types; implementations, however, already have representation on the committee, and will veto any truly unimplementable things in this space.

4.1.4 ORDER(X, Y) already has public-facing API implications

The “normalized” versions of types will inevitably show up in function signatures; if different compilers on the same platform produced different orderings, compilation units from different compilers would refuse to link, despite both “signatures” beings spelled identically.

Letting the compiler vendors synchronize based on the mangling scheme or something equivalently useful is the same forum as the current ABI discussion forums; it seems the appropriate venue to standardize the ordering anyway, without making WG21 duplicate the work.

4.2 Desirable properties of ORDER(x, y)

4.2.1 Stability

The order should be the same across compilation units. This is key for generating ABI-compatible vtables, for instance.

It also should be stable through time. An order based on an ABI-mangling scheme satisfies this notion, at least.

4.2.2 Free-standing

The order should be available on freestanding implementations. One crucial use-case is replacing the exception mechanism on those platforms with return values of `std::expected` or similar, and compromising that is undesirable.

4.2.3 Consistency with `type_info::before()`

Ideally, whenever `typeid(x).before(typeid(y)) == true`, then `ORDER(x, y) == std::strong_ordering::less`. The converse obviously cannot be true, since `ORDER(x, y)` is finer than the order induced by `type_info::before()`.

However, the standard currently says

The names, encoding rule, and collating sequence for types are all unspecified and may differ between programs.

Since this paper requires that the order not differ between programs, exposing this as a normative requirement is impossible without tightening this wording, which is outside the scope of this paper.

4.2.4 Self-consistency

The ordering should be self-consistent, that is, for all possible template arguments `T`, `U`, and any unary template `some_template`:

```
ORDER(T, U) == ORDER(some_template<T>, some_template<U>).
```

Implementations are encouraged to satisfy this principle, but are not required to.

4.2.5 Reflection compatibility

Any `operator<=>(std::meta::info, std::meta::info)` should be consistent with this one.

While `std::meta::info` [P2320R0] objects can reflect more entities than just types and values (mainly: expressions), any ordering defined on them should be finer than this one, and specifically consistent with it. We should do something that reflection can subsume.

However, it doesn't seem like that proposal will define an ordering on `info` objects, and we need this solved sooner than that.

5 Proposal

5.1 Semantics

Regardless of the syntax chosen, the semantics would be the following.

Let `X` and `Y` be (possibly `cv-ref`) qualified types.

`ORDER(X, Y)` is an constant expression of type `std::strong_order`.

- `std::same_as<X, Y> == true` if and only if `ORDER(X, Y) == std::strong_ordering::equal`.
- Otherwise, `ORDER(X, Y)` is either `std::strong_ordering::less` or `std::strong_ordering::greater`. Which of those is implementation-defined, subject to the following semantic constraints.

Implementations must define `ORDER(X, Y)` such that it is an ordering, that is, it is transitive and antisymmetric; that is

- if `ORDER(X, Y) == std::strong_ordering::less`, then `ORDER(Y, X) == std::strong_ordering::greater` and vice versa
- if `ORDER(X, Y) == std::strong_ordering::less` and `ORDER(Y, Z) == std::strong_ordering::less`, then `ORDER(X, Z)` is also `std::strong_ordering::less`.

It is implementation-defined whether the order `ORDER(X, Y)` is finer than the one implied by `std::type_info::before`. That is,

- if `typeid(X).before(typeid(Y))`, then `ORDER(X, Y) == std::strong_ordering::less`.

Note: the converse is not possible - `ORDER(X, Y)` does not strip *cv-ref* qualifiers.

Please also note that `X` and `Y` are types, and therefore `typeid` does not have dynamic typing semantics in this case.

Note: Implementations are encouraged to do this if possible, and making it “implementation-defined” places a requirement on implementations to document whether they did so.

5.2 Syntax

Any syntax will do for the use-case, but some better than others. This section explores the trade-offs in syntax choice.

The authors recommend the first option.

5.2.1 Option 1 (preferred): a variable template `std::type_order_v<T, U>`

Specifically:

```
template <typename T, typename U>
inline constexpr std::strong_ordering type_order_v = ORDER(T, U); /* see below */
template <typename T, typename U>
struct type_order : integral_constant<strong_ordering, type_order_v<T, U>> {};

// as a separate library proposal, once member packs make it
template <typename... Ts>
using ...typemultiset = /* pack of Ts, sorted by type_order_v */;
template <typename... Ts>
using ...typeset = /* unique ...typemultiset<Ts...>... */;
```


This seems like a pretty good choice. It does not need a new keyword, only depends on `<compare>`, and the name seems relatively discoverable.

We could, in fact, put it into `<compare>` if we didn't want `<type_traits>` to depend on it.

It's also freestanding, since it doesn't depend on `<typeinfo>`.

We should also allow, as a special provision, the arguments to the above metafunctions to be incomplete.

5.2.2 Option 2: variable template `std::entity_ordering<X, Y>`

Specifically:

```
template <universal template X, universal template Y>
inline constexpr std::strong_order entity_order_v = ORDER(X, Y); /* see below */
template <universal template X, universal template Y>
struct entity_order : integral_constant<strong_ordering, entity_order_v<X, Y>> {};
```

This is a better option than Option 1 if we get universal template parameters, as we really want to also order class templates, not just types.

However, without universal template parameters, we really don't have much of a choice but to reach for Option 1.

The name `entity_order` is also slightly less obvious than `type_order`, but metaprogrammers shouldn't have trouble finding either.

5.2.3 Option 3: reflection

Specifically:

```
constexpr std::partial_order entity_order(std::meta::info x, std::meta::info y) {
    return __comparable(x, y) ? ORDER(x, y) : std::partial_order::unordered;
}
```

We could standardize a type order as a function on `std::meta::info` objects in `std::meta`. However, once we're in `std::meta::info` space, it's more difficult to know which reflections are comparable and which aren't, so such a function would need to return a `std::partial_order`, which seems decidedly less desirable.

It also means we'd need to pass the correct kind of reflection into the ordering function, which is a bit less intuitive than just `decltype` or just the template parameter that we already have.

5.2.4 Option 4: heterogeneous constexpr `std::type_identity::operator<=>`

Specifically:

```
template <typename T, typename U>
constexpr std::strong_order operator<=>(std::type_identity<T>, std::type_identity<U>);
```

Pros:

- No new names.

Cons:

- Less discoverable than a new type-trait
- Requires template instantiations of `type_identity<T>`, `type_identity<U>`, as well as `operator<=>` overload resolution and substitution, which is quite expensive compared to a single `type_order_v<T, U>` direct substitution and lookup.
- An extra `operator<=>` overload in the `std` namespace is a drag on overload resolution
- adds `<compare>` to `<type_traits>`, since that's where `type_identity` is
- too cute?

- doesn't work for nontypes

5.2.5 Option 5: `constexpr std::__lift<arg>::operator<=>`

This option means we add `template <universal template> struct __lift {};` into `<type_traits>` and define `operator<=>` for it.

Pros:

- ... we'll need a lift sooner or later?

Cons:

- all the cons of `type_identity`
- even more nonobvious
- still needs a new name
- needs a tutorial to find

5.2.6 Non-Option: `constexpr bool std::type_info::before()`

It would be nice, but alas, operates on *cv-unqualified* versions of the referenced type, so it's not sufficient.

`constexpr std::strong_order(std::type_info, std::type_info)` has similar issues.

6 FAQ

6.1 Why should this be standardized?

Because we have no way to reliably order types across compilation units at compile-time.

6.2 Why not wait for reflection?

It's a good question. However, reflection will do *nothing* for this problem by itself; the user will still have to implement ordering using `constexpr` functions, which have no hope of being as fast as a compiler-provided built-in.

User-programmed functions also won't adapt to language evolution; this feature will.

Finally, sorting is arbitrary; having it be consistent throughout the software ecosystem is potentially a great enabler of interoperability.

6.3 But couldn't this be done faster with reflection?

No; Peter Dimov shares an interesting anecdote.

I have in Mp11 the algorithm `mp_unique`, which takes a list of types and removes the duplicates. In the course of writing the reflection papers, their authors occasionally took Mp11 code examples and tried to show how they are elegantly implemented using value-based reflection metaprogramming.

So, you take a `vector<info>` that contains types, and then you simply apply the existing algorithm `std::unique` to it, et voila... oh wait.

`std::unique` wants a sorted range, and you can't `std::sort` the info vector, because info objects aren't ordered, even when they refer to types.

7 Acknowledgements

Thanks to all of the following:

- Davis Herring for his suggestions on ordering non-type template parameters.

- Ville Voutilainen for his critique of examples, and providing a simple way of explaining the motivation
- Peter Dimov for a helpful anecdote, now in the FAQ.
- Erich Keane for pushing us back to the “implementation-defined” territory.

8 Proposed Wording

In [compare.syn], add

```
template <class T, class U>
struct type_order : integral_constant<strong_ordering, TYPE_ORDER(T, U)> {};
template <class T, class U>
inline constexpr strong_ordering type_order_v = type_order::value;
```

At the end of [cmp], just before [support.coroutine], add:

17.11.7: Type Ordering

- ¹ Define the constant expression `TYPE_ORDER(X, Y)` for arbitrary (possibly *cv*-qualified or reference) types `X` and `Y` as follows - if `is_same_v<X, Y>`, then `TYPE_ORDER(X, Y)` is `strong_ordering::equal`
- (1.2) - otherwise, it is implementation-defined whether `TYPE_ORDER(X, Y)` is `strong_ordering::less` or `strong_ordering::greater`, subject to the following constraints - antisymmetry: `TYPE_ORDER(X, Y)` is `strong_ordering::less` if and only if `TYPE_ORDER(Y, X)` is `strong_ordering::greater` - transitivity: if both `TYPE_ORDER(X, Y)` and `TYPE_ORDER(Y, Z)` are `strong_ordering::less`, then `TYPE_ORDER(X, Z)` is also `strong_ordering::less`. The name `type_order` denotes a *Cpp17BinaryTypeTrait* (20.15.2) with a base characteristic of `integral_constant<strong_ordering, TYPE_ORDER(X, Y)>`.
- ³ *Recommended practice:* `[_TYPE_ORDER?](X, Y)_@` should be a finer order than the one induced by `type_info::before` if such an implementation is possible on a given platform.

9 Full order proposal (former alternative)

This is a poor first attempt at figuring out what it would take to specify this in the C++ standard. It seems like “implementation-defined” is a better approach.

9.1 Approach

1. We define a **lowering** to a *key-tuple* for every entity in the language.
2. The order is then defined on these *key-tuples*.

9.2 Structure of key-tuples

Every *key-tuple* is of the form `(element...)`.

where an element is one of:

- *atom* (see [atoms](#))
- *key-tuple*

These tuples are then ordered lexicographically (ties broken in favor of shorter tuple), atoms before tuples.

Let us name this transformation as `sort_key(entity)`.

The rest of the paper is concerned with defining this transformation.

9.3 Named Scopes

A type is ordered by appending `sort_key(...)` to the named scope it is declared in. The following are *named scopes*:

1. namespaces
2. classes
3. functions
4. lambdas
5. concepts

Starting with the global namespace, `sort_key(global) = ()`. Any type `T` declared in the global namespace shall have a defined `sort_key` operation that resolves to `(sort_key(T))`.

9.3.1 Example 1: class `foo` is declared in struct `bar`:

```
struct bar { class foo; }
```

```
sort_key(foo) = (sort_key(bar), sort_key(foo)) = ((type, bar), (type, foo, ))
```

This shall hold for any of the above named scopes.

9.3.2 Example:

Given

```
namespace foo::bar {
    struct i;
}

namespace baz {
    struct j;
}
```

Then:

- `sort_key(foo::bar::i)` is `((namespace, foo), (namespace, bar), (type, i,))`.
- `sort_key(baz::j)` is `((namespace, baz), (type, j,))`

When compared, the result is that `baz::j < foo::bar::i`, since namespace `baz` precedes namespace `foo`.

9.4 Atoms

The atoms of *key-tuples* are ordered as follows:

1. kinds (see [kinds](#))
2. simple names (including empty string) (see [names](#))
3. qualifiers (see [qualifiers](#))
4. `[]` (array of unknown bound)
5. `[n]` (array of known bound `n`) (ordered by `n` internally)
6. `*` (pointer)
7. ellipsis (`...` in `f(...)`)
8. parameter pack (`...` in `typename...`)

9.5 Kinds

There are the following kind tokens that can appear in *key-tuples*.

1. value
2. namespace
3. type
4. class template
5. type alias template

6. variable template
7. concept
8. function

Note: everything but “values” is pretty simple, but we haven’t dealt with values extensively yet with the R1 of this paper, though we should just defer to \Leftrightarrow and require a default strong structural ordering for values that may be template arguments.

9.5.1 Identifiers

9.5.1.1 Simple Names

Most names are strings that are valid (atomic) identifiers. Those are just themselves:

```
namespace foo::bar { struct baz; }
```

`foo`, `bar` and `baz` are such atomic identifiers.

9.5.1.2 Anonymous Namespace

Anonymous namespaces shall be represented with the `!` character, as it cannot be represented by the empty string and cannot collide with any user defined names;

Example:

```
namespace a { namespace { struct s; } }

sort_key(a::s) = ((namespace, a), (namespace, "!"), (type, s, ))
```

9.5.1.3 Unnamed entities

Unnamed entities are all given a name that is not an identifier (but is, in fact, a tuple), and are then numbered consecutively, starting with zero, based on their name-scope.

Name-scopes are namespaces, classes, unions, functions, and enumerations.

Function declarations are name-scoped to the function itself.

Consider a lambda that appears as a default argument of a function template:

```
template <typename T>
void f(T x = []{ return T{0}; }());
//           ~~~~~~ this one
```

The *key-tuple* for `f<int>` is:

```
(function, (f, (type, int)), (type, void), ((type, int)))
```

The *key-tuple* for the lambda is:

```
((function, (f, (type, int)), (type, void), ((type, int))), (type, (lambda, 0), )).
```

Note: because of the regular structure of *key-tuples*, such anonymous classes will compare greater than any entity that has a simple identifier, due to tuples comparing greater than atoms (which simple names are).

9.5.1.3.1 Lambda types

Types of lambda objects are ordered first by where they are declared, then by declaration order.

In effect, we assign them the name `(lambda, #)` where `#` is the count of other unnamed entities in the name scope.

```

namespace Banana {
    auto i = [] (int) -> void {}; // 0th lambda instantiated in Banana
}

namespace Apple {
    auto i = [] (float) -> int {}; // 0th lambda instantiated in Apple
    auto j = [] () -> std::string {}; // 1st lambda instantiated in Apple
}

```

These would produce the following tuples:

```

sort_key(decltype(Banana::i)) = ((namespace, Banana), (type, (lambda, 0), ));
sort_key(decltype(Apple::i)) = ((namespace, Apple), (type, (lambda, 0), ));
sort_key(decltype(Apple::j)) = ((namespace, Apple), (type, (lambda, 1), ));

```

Note: the empty bit after the identifier is the empty qualifier pack.

9.5.1.3.2 Unnamed struct and union types

They are named, respectively, (class, #) and (union, #).

9.5.2 Namespaces

The `sort_key(namespace-name)` is (namespace, identifier).

This means that namespaces are ordered alphabetically by comparing namespace names at the same rank. A namespace comes before any of its subnamespaces.

Example:

```

namespace outer1 {
    struct i;
}

namespace outer2 {
    namespace inner1 {
        struct i;
    }
    namespace inner2 {
        struct i;
    }
}

```

The order of the three structs w/ type i types shall be

```
sort_key(outer1::i) < sort_key(outer2::inner1::i) < sort_key(outer2::inner2::i).
```

9.5.3 Types

The `sort_key` of a type is (type, <identifier>, <qualifiers>).

The <identifier> bit is a bit complicated, so let's deal with the qualifiers first.

Note: any name-scopes the `type` is declared in are part of the parent *key-tuple*. The `identifier` portion is complicated because of possible template arguments for types that are template specializations.

9.5.3.1 Qualifiers

Qualifiers are each assigned a score

```
&: 1
&&: 2
const: 3
volatile: 6
```

and ordering lowest-first after summing them.

Therefore, for an unqualified type `T`, the order of all possible qualified types would be:

```
0 T
1 T &
2 T &&
3 T const
4 T const &
5 T const &&
6 T volatile
7 T volatile &
8 T volatile &&
9 T const volatile
10 T const volatile &
11 T const volatile &&
```

The remainder of the paper concerns itself only with unqualified types.

9.5.4 Ordering Scalar Types

All scalar types are built-in types, except for enumerations, which should be ordered according to their namespaced names.

Unfortunately, some of the built-in types do not have names, only type aliases (such as `decltype(nullptr)`).

The intention is for built-in scalar types to be ordered before any compound types.

Built-in types with simple names should be ordered before any types that reference other types.

In particular, scalar types shall be ordered as follows:

1. `void` comes first because it's not reifiable,
2. the type of `std::nullptr_t` as the first monostate
3. any other monostates, if added, sorted alphabetically by their common names (to be specified explicitly if added)
4. `bool` as the first bi-state
5. any other bi-states, if added, sorted alphabetically.
6. Raw-memory types (`char`, `signed char`, `unsigned char`) (`std::byte` is an enumeration in `std` so it falls under different rules)
7. Integral types in order of size, signed before unsigned (`short`, `unsigned short`, `int`, `unsigned int`, `long`, `unsigned long`, `long long`, `unsigned long long`, followed by any implementation-defined wider integral types like `__int128_t` etc.). Intersperse any implementation-defined built-in integral types as needed between the above.
8. Any remaining character types that are not type-aliases of any of the above, including unicode, according to the following rules: smallest first, unicode-specific variants after non-unicode variants.
9. Floating-point types, in order of size. In case of ties, `float`, `double` and `long double` come before any floating point types.
10. Function types (internally ordered by rules in section [Function Types](#))
11. Pointer types (internally ordered by their pointee-type)
12. Pointer-to-member types (internally ordered by pointee-type)

Class types shall be ordered according to the rules below, see [Ordering Compound Types]

9.5.5 Ordering Array Types

Array types shall be ordered after scalar types but before class types.

The `sort_key(T[]) = ([], sort_key(T))` and the `sort_key(T[n]) = ([n], sort_key(T))`.

The intention is to order arrays first internally by element type, then by rank, then by rank bounds, lowest first. Arrays of unknown bounds come before arrays of known bounds.

So the order of the following, for a given type `T`:

```
T[]
T[10]
T[11]
T[] [2]
T[10] [2]
T[3] [2]
```

shall be ordered `T[] < T[10] < T[11] < T[] [2] < T[3] [2] < T[10] [2]`, and

`sort_key(T[0]) = (type, ([], (type, T,)))`

`sort_key(T[10][2]) = (type, ([2], sort_key(T[10]))) = (type, ([2], (type, ([10], (type, T,)))))`

9.5.6 Ordering Class Types

9.5.6.1 Ordering Simple Class Types

Class types shall be greater than scalar types.

Since we cannot redeclare two types with the same name, class types shall be ordered alphabetically.

```
struct Apple {};
class Banana {};
struct Carrot {};
```

Would be ordered as `Apple < Banana < Carrot`

As such, we define sort key as:

`sort_key(Apple) = (type, Apple,)`

`sort_key(Banana) = (type, Banana,)`

`sort_key(Carrot) = (type, Carrot,)`

9.5.6.2 Non Type Template Parameters

NTTPs shall first be ordered by their type, then their value.

Given:

```
template <auto T>
struct s {
    decltype(T) i = T;
};

s<1u> a;
s<1.0f> b;
```

`sort_key(s<1u>) = ((type, (s, sort_key(1u))))`

We can define `sort_key` of `1u` as: `sort_key(1u) = (sort_key(decltype(1u)), 1)`

`s<1u>` shall be ordered before `s<1.0f>`, as integral types come before floating point types.

NTPs of the same type shall be lexicographically ordered by their scalar subobjects. Meaning

```
struct F final {
    struct G final {
        int h;
        int i;
    } g;
    int j;
};
```

```
F f{{0,1}, 2};
F f2{{1,2}, 3};
```

`sort_key(s<f>) < sort_key(s<f2>);`

NTPs of the same pointer or reference type shall be ordered by instantiation order.

9.5.6.3 Ordering Class Template Specializations

Class templates shall be ordered by:

- 1) Class name, alphabetically.
- 2) Template arguments, applied lexicographically.

For example, given:

```
template <typename T, typename U>
struct Apple;

struct Banana;
struct Carrot;

Apple<Banana, Carrot>;
Apple<Banana, Banana>;
Apple<Carrot, Carrot>;
```

Note, `sort_key(<parameter>)...` will be used to denote a tuple where `sort_key` has been applied to all parameters.

For `void f(Foo, Bar) sort_key(<parameter>)...` would mean `(sort_key(Foo), sort_key(Bar))`

`sort_key` of a class template shall be defined as:

`sort_key(<class template>) = (type, (<name>, (sort_key(<parameter>)...)))`

So

`sort_key(Apple<Banana, Carrot>) = (type, (Apple, (sort_key(Banana), sort_key(Carrot)),)`

`sort_key(Apple<Banana, Carrot>) = (type, (Apple, ((type, Banana,), (type, Carrot,)),)`

Note: the empty bit after the identifier is the empty qualifier pack.

The above would be ordered `sort_key(Apple<Banana, Banana>)`, `sort_key(Apple<Banana, Carrot>)`, `sort_key(Apple<Carrot, Carrot>)`.

9.5.6.4 Function Types

Function types shall be ordered by

1. Return type
2. Parameters, lexicographically.

The `sort_key` of a function shall be defined as:

```
sort_key(<function>) = (function, <name>, sort_key(<return type>), (sort_key(<parameter>)...))  
void foo(int i);
```

This function can be represented by: (function, (foo, (type, void), ((type, int))))

```
void foo(int)  
void foo(int, double)
```

```
sort_key(void foo(int)) = (function, foo, (type, void), ((type, int)))
```

```
sort_key(void foo(int, double)) = (function, foo, (type, void), ((type, int), (type, double)))
```

So, the type of `void foo(int)` would precede the type of `void foo(int, double)`

9.5.6.5 Member Function Types

Function types shall be ordered by

1. Return type
2. The type of the class it is a member of.
3. Parameters, lexicographically.

The sort key of a member function shall be defined as:

```
sort_key(<member function>) =  
(function, (<name>, sort_key(<class>)), sort_key(<return type>), (sort_key(<parameter>)...)))  
struct Foo {  
    void bar(int i, float j);  
};
```

```
sort_key(Foo::bar) =
```

```
(type, Foo, ), (function, (bar, (type, Foo, )), (type, void), ((type, int, ), (type, float, ))))
```

9.5.6.6 Variadic Function Types

Variadic function shall be ordered in a similar way. In a variadic function, the last argument is a variadic argument. A variadic argument shall be ordered immediately after its underlying type.

Given:

```
void foo(Foo);  
void foo(Foo...);
```

In this case, the type of `void foo(Foo...)` is ordered immediately after the type of `void foo(Foo)`.

We can represent these as:

```
(function (type, void) (type, Foo, ))  
(function (type, void) (type, Foo, ...))
```

9.5.6.7 Parameter Packs

Parameter are ordered as class templates.

Given:

```
template<class... Types>
struct Tuple {};

class Foo {};
class Bar {};

Tuple<> t0;
Tuple<int> t1;
Tuple<Foo> t2;
Tuple<Bar> t3;
Tuple<Foo, Bar> t4;
```

would be ordered: Tuple<> < Tuple<int> < Tuple<Bar> < Tuple<Foo> < Tuple<Foo, Bar>

9.5.6.8 Ordering Class Templates

Kinds of templates are ordered first by name, then by template arguments.

Given:

```
template <template <template<typename> class> class Template>
struct two{};

template <template <typename> class> struct one{};

template <typename> struct zero{};

zero<int> value0;
one<zero> value1;
two<one> value2;
```

These are represented by tuples:

```
sort_key(zero<int>) = (type, (zero, (type, int)))
sort_key(one<zero>) = (type, (one, (class_template, zero)))
sort_key(two<one>) = (type, (two, (class_template, one)))
```

9.5.6.9 Variable Templates

Variable templates are ordered by name, then by template parameter.

```
sort_key(<variable_template>) = (variable_template, (<name>, (sort_key(<template_parameter>)...)))
template <typename F, typename S>
constexpr std::pair<F, S> pair_one_two = {1, 2};
```

the type of pair_one_two<int, double> can be represented as:

```
sort_key(pair_one_two<int, double>) = (variable_template, (pair_one_two, (type, int), (type, double)))
```

9.5.6.10 Alias Templates

Alias templates are ordered alphabetically by name.

```
sort_key(<alias_template>) = (alias_template, <name>)
```

Given

```
template< class T >
using remove_cvref_t = typename remove_cvref<T>::type;

sort_key(remove_cvref_t) = (alias_template, remove_cvref_t)
```

9.5.6.11 Concepts

Concepts are ordered in a similar manner to variable templates.

```
sort_key(<concept>) = (concept, (<name>, (sort_key(<template_parameter>)...)))

template <typename T, typename F = decltype([](T){})>
concept f = requires (T i, F f = [](T){}) {
    {f(i)} -> std::convertible_to<void>;
};
```

In order to order the type of the lambda declared in concept `f`, concept `f` must be comparable with other types.

Concepts shall be ordered first by name, then by template arguments.

```
sort_key(f<int>) = (concept, (f, (type, int), (lambda, 0)))
```

10 Appendix A: building apply_canonicalized

We will need a small metaprogramming library; a filter is difficult to do otherwise.

```
struct undefined;
template <typename... Ts> struct list {};

// apply<F, list<Ts...>> -> F<Ts...>
template <template <typename...> typename, typename> extern undefined _apply;
template <template <typename...> typename F, template <typename...> typename L,
        typename... Ts>
F<Ts...> _apply<F, L<Ts...>>;
template <template <typename...> typename F, typename List>
using apply = decltype(_apply<F, List>);

// concatenate<list<Ts...>, list<Us...>, list<Vs...>> -> list<Ts..., Us..., Vs...>
template <typename...> extern undefined _concatenate;
template <typename... Ts> list<Ts...> _concatenate<list<Ts...>>;
template <typename... Ts, typename... Us, typename... Lists>
decltype(_concatenate<list<Ts..., Us..., Lists...>())
_concatenate<list<Ts...>, list<Us...>, Lists...>;
template <typename... Ts>
using concatenate = decltype(_concatenate<Ts...>);

// select: list<T> if true, list<> if false
template <bool v, typename T> extern list<> _select;
template <typename T> list<T> _select<true, T>;

template <bool v, typename T>
using select = decltype(_select<v, T>);
```

Canonicalization is now just a basic not-in-place quicksort-ish thing:

```

template <typename.../*empty*/> extern list<> _canon;
template <typename... Ts>
using canonicalized = decltype(_canon<Ts...>);

// a canonicalized T is just T
template <typename T>
list<T> _canon<T>;

template <typename T, typename... Ts>
concatenate<
    // canonicalized things less than T
    apply<canonicalized, concatenate<select<(ORDER(Ts, T) < 0), Ts>...>>,
    list<T> /*T*/, // ~~~~~~
    // canonicalized things greater than T
    apply<canonicalized, concatenate<select<(ORDER(Ts, T) > 0), Ts>... >>
    > // ~~~~~~
_canon<T, Ts...>;

```

We now have `canonicalized<Ts...>` - but this still leaves `list` as a special type which we'd rather not expose to the user. Onto `apply_canonicalized`:

```

template <template <typename...> typename F, typename... Ts>
using apply_canonicalized = apply<F, canonicalized<Ts...>>;

```

10.1 Full code listing as tested and implemented

Here for completeness, feel free to skip.

```

#include <compare>
#include <type_traits>

struct undefined;

#define ORDER(x, y) type_order_v<x, y>

// in <type_traits>
template <typename X, typename Y>
constexpr inline std::strong_ordering type_order_v;

template <template <typename...> typename, typename>
extern undefined _apply;

template <template <typename...> typename F, template <typename...> typename L,
        typename... Ts>
F<Ts...> _apply<F, L<Ts...>>;

template <template <typename...> typename F, typename List>
using apply = decltype(_apply<F, List>);

// some user-type
template <auto x>
struct value_t : std::integral_constant<decltype(x), x> {};
template <auto x>
inline constexpr value_t<x> value_v{};

```

```

// built-in
template <auto x, auto y>
constexpr inline std::strong_ordering type_order_v<value_t<x>, value_t<y>> =
    x <=> y;

template <typename... Ts>
struct list {};

template <typename...>
extern undefined _concatenate;
template <typename... Ts>
list<Ts...> _concatenate<list<Ts...>>;
template <typename... Ts, typename... Us, typename... Lists>
decltype(_concatenate<list<Ts..., Us...>, Lists...>)
    _concatenate<list<Ts...>, list<Us...>, Lists...>;

template <typename... Ts>
using concatenate = decltype(_concatenate<Ts...>);

template <bool v, typename T>
extern list<> _select;
template <typename T>
list<T> _select<true, T>;

template <bool v, typename T>
using select = decltype(_select<v, T>);

template <typename...>
extern list<> _canon;
template <typename... Ts>
using canonicalized = decltype(_canon<Ts...>);

template <typename T>
list<T> _canon<T>;

template <typename T, typename... Ts>
concatenate<
    apply<canonicalized, concatenate<select<<ORDER(Ts, T) < 0>, Ts>...>>,
        list<T>,
        apply<canonicalized, concatenate<select<<ORDER(Ts, T) > 0>, Ts>...>>
    >
    _canon<T, Ts...>;

static_assert(std::same_as<canonicalized<value_t<0>, value_t<-1>, value_t<-1>, value_t<1>>, list<value_t<0>, value_t<-1>, value_t<-1>, value_t<1>>>, list<value_t<0>, value_t<-1>, value_t<-1>, value_t<1>>>);

```

11 References

- [P1985R3] Gašper Ažman, Mateusz Pusz, Colin MacLean, Bengt Gustafsonn, Corentin Jabot. 2022-09-17. Universal template parameters. <https://wg21.link/p1985r3>
- [P2300R7] Eric Niebler, Michał Dominiak, Georgy Evtushenko, Lewis Baker, Lucian Radu Teodorescu, Lee

Howes, Kirk Shoop, Michael Garland, Bryce Adelstein Lelbach. 2023-04-21. ‘std::execution’.
<https://wg21.link/p2300r7>

[P2320R0] Andrew Sutton, Wyatt Childers, Daveed Vandevoorde. 2021-02-15. The Syntax of Static Reflection.
<https://wg21.link/p2320r0>

[P2841R1] Corentin Jabot, Gašper Ažman. 2023-10-14. Concept Template Parameters.
<https://wg21.link/p2841r1>